

Final Project

June 4, 2022

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.formula.api as sm
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.metrics import
    ↳ mean_squared_error, r2_score, roc_curve, auc, precision_recall_curve, auc,
    ↳ make_scorer, recall_score, accuracy_score, precision_score
from sklearn.model_selection import KFold
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV, ParameterGrid
from sklearn.ensemble import
    ↳ BaggingRegressor, BaggingClassifier, RandomForestRegressor, RandomForestClassifier, StackingClassifier
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import Ridge, RidgeCV, Lasso, LassoCV
import itertools as it
from sklearn.model_selection import StratifiedKFold, KFold

from sklearn.tree import export_graphviz
import xgboost as xgb
from six import StringIO
from IPython.display import Image
import pydotplus
import time as time
```

0.1 Data Cleaning

```
[2]: data = pd.read_csv('mushrooms.csv')

[3]: data.rename(columns = {'class': 'classes'}, inplace=True)

[4]: data.columns.to_list()
```

```
[4]: ['classes',
      'cap-shape',
      'cap-surface',
      'cap-color',
      'bruises',
      'odor',
      'gill-attachment',
      'gill-spacing',
      'gill-size',
      'gill-color',
      'stalk-shape',
      'stalk-root',
      'stalk-surface-above-ring',
      'stalk-surface-below-ring',
      'stalk-color-above-ring',
      'stalk-color-below-ring',
      'veil-type',
      'veil-color',
      'ring-number',
      'ring-type',
      'spore-print-color',
      'population',
      'habitat']
```

```
[5]: data.columns = data.columns.str.strip().str.lower().str.replace('-', '_')
```

```
[6]: data.isna().sum()
```

```
[6]: classes          0
     cap_shape        0
     cap_surface      0
     cap_color        0
     bruises          0
     odor             0
     gill_attachment  0
     gill_spacing     0
     gill_size        0
     gill_color       0
     stalk_shape      0
     stalk_root       0
     stalk_surface_above_ring  0
     stalk_surface_below_ring  0
     stalk_color_above_ring  0
     stalk_color_below_ring  0
     veil_type        0
     veil_color       0
     ring_number      0
```

```

ring_type          0
spore_print_color  0
population         0
habitat           0
dtype: int64

```

```
[7]: data.head()
```

```

[7]:  classes cap_shape cap_surface cap_color bruises odor gill_attachment \
0      p      x      s      n      t      p      f
1      e      x      s      y      t      a      f
2      e      b      s      w      t      l      f
3      p      x      y      w      t      p      f
4      e      x      s      g      f      n      f

      gill_spacing gill_size gill_color  ... stalk_surface_below_ring \
0      c      n      k  ...      s
1      c      b      k  ...      s
2      c      b      n  ...      s
3      c      n      n  ...      s
4      w      b      k  ...      s

      stalk_color_above_ring stalk_color_below_ring veil_type veil_color \
0      w      w      p      w
1      w      w      p      w
2      w      w      p      w
3      w      w      p      w
4      w      w      p      w

      ring_number ring_type spore_print_color population habitat
0      o      p      k      s      u
1      o      p      n      n      g
2      o      p      n      n      m
3      o      p      k      s      u
4      o      e      n      a      g

[5 rows x 23 columns]

```

```

[8]: classes = {
      'e': 'edible',
      'p': 'poisonous'
    }

    cap_shapes = {
      'b': 'bell',
      'c': 'conical',
      'x': 'convex',

```

```

    'f': 'flat',
    'k': 'knobbed',
    's': 'sunken'
}

cap_surfaces = {
    'f': 'fibrous',
    'g': 'grooves',
    'y': 'scaly',
    's': 'smooth'
}

cap_colors = {
    'n': 'brown',
    'b': 'buff',
    'c': 'cinnamon',
    'g': 'gray',
    'r': 'green',
    'p': 'pink',
    'u': 'purple',
    'e': 'red',
    'w': 'white',
    'y': 'yellow'
}

bruise_class = {
    't': 'bruises',
    'f': 'no_bruises'
}

odors = {
    'a': 'almond',
    'l': 'anise',
    'c': 'creosote',
    'y': 'fishy',
    'f': 'foul',
    'm': 'musty',
    'n': 'none',
    'p': 'pungent',
    's': 'spicy'
}

gill_attachments = {
    'a': 'attached',
    'd': 'descending',
    'f': 'free',
    'n': 'notched'
}

```

```

}

gill_spacings = {
    'c': 'close',
    'w': 'crowded',
    'd': 'distant'
}

gill_sizes = {
    'b': 'broad',
    'n': 'narrow'
}

gill_colors = {
    'k': 'black',
    'n': 'brown',
    'b': 'buff',
    'h': 'chocolate',
    'g': 'gray',
    'r': 'green',
    'o': 'orange',
    'p': 'pink',
    'u': 'purple',
    'e': 'red',
    'w': 'white',
    'y': 'yellow'
}

stalk_shapes = {
    'e': 'enlarging',
    't': 'tapering'
}

stalk_roots = {
    'b': 'bulbous',
    'c': 'club',
    'u': 'cup',
    'e': 'equal',
    'z': 'rhizomorphs',
    'r': 'rooted',
    '?': 'NA'
}

stalk_surface_above_rings = {
    'f': 'fibrous',
    'y': 'scaly',
    'k': 'silky',

```

```

        's': 'smooth'
    }

    stalk_surface_below_rings = {
        'f': 'fibrous',
        'y': 'scaly',
        'k': 'silky',
        's': 'smooth'
    }

    stalk_color_above_rings = {
        'n': 'brown',
        'b': 'buff',
        'c': 'cinnamon',
        'g': 'gray',
        'o': 'orange',
        'p': 'pink',
        'e': 'red',
        'w': 'white',
        'y': 'yellow'
    }

    stalk_color_below_rings = {
        'n': 'brown',
        'b': 'buff',
        'c': 'cinnamon',
        'g': 'gray',
        'o': 'orange',
        'p': 'pink',
        'e': 'red',
        'w': 'white',
        'y': 'yellow'
    }

    veil_types = {
        'p': 'partial',
        'u': 'universal'
    }

    veil_colors = {
        'n': 'brown',
        'o': 'orange',
        'w': 'white',
        'y': 'yellow'
    }

    ring_numbers = {

```

```

    'n': 'none',
    'o': 'one',
    't': 'two'
}

ring_types = {
    'c': 'cobwebby',
    'e': 'evanescent',
    'f': 'flaring',
    'l': 'large',
    'n': 'none',
    'p': 'pendant',
    's': 'sheathing',
    'z': 'zone'
}

spore_print_colors = {
    'k': 'black',
    'n': 'brown',
    'b': 'buff',
    'h': 'chocolate',
    'r': 'green',
    'o': 'orange',
    'u': 'purple',
    'w': 'white',
    'y': 'yellow'
}

populations = {
    'a': 'abundant',
    'c': 'clustered',
    'n': 'numerous',
    's': 'scattered',
    'v': 'several',
    'y': 'solitary'
}

habitats = {
    'g': 'grasses',
    'l': 'leaves',
    'm': 'meadows',
    'p': 'paths',
    'u': 'urban',
    'w': 'waste',
    'd': 'woods'
}

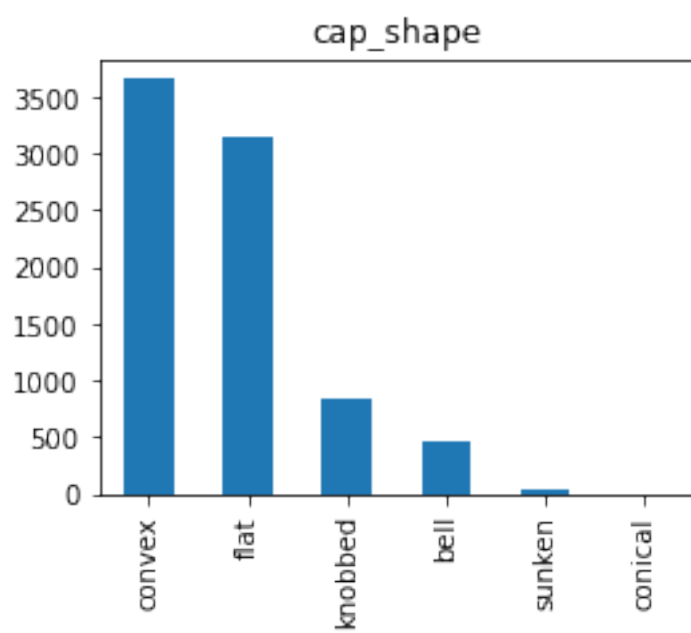
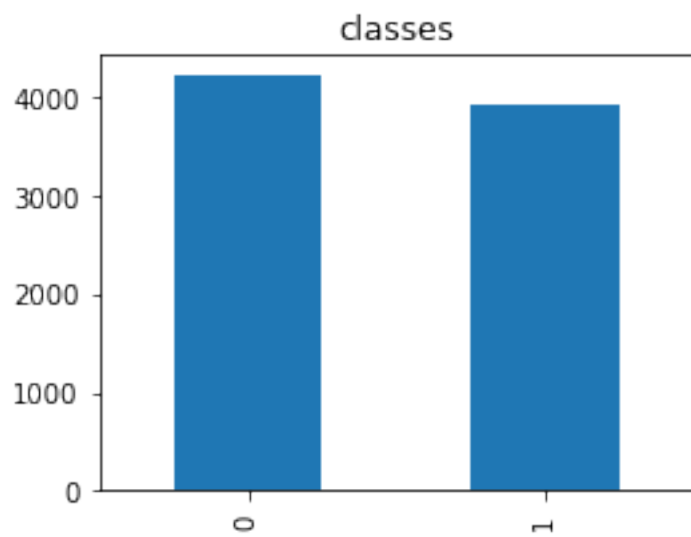
```


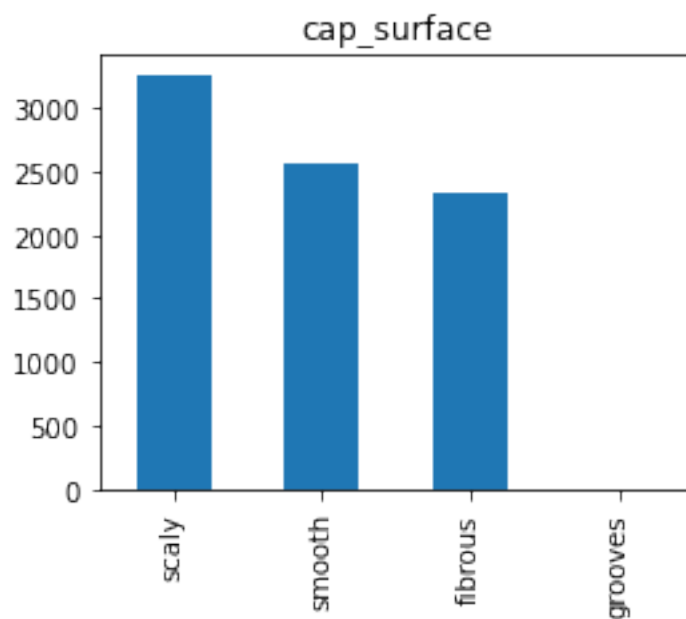
```
[9]: data.replace({'classes': classes,
                  'cap_shape': cap_shapes,
                  'cap_surface': cap_surfaces,
                  'cap_color': cap_colors,
                  'bruises': bruise_class,
                  'odor': odors,
                  'gill_attachment': gill_attachments,
                  'gill_spacing': gill_spacings,
                  'gill_size': gill_sizes,
                  'gill_color': gill_colors,
                  'stalk_shape': stalk_shapes,
                  'stalk_root': stalk_roots,
                  'stalk_surface_above_ring': stalk_surface_above_rings,
                  'stalk_surface_below_ring': stalk_surface_below_rings,
                  'stalk_color_above_ring': stalk_color_above_rings,
                  'stalk_color_below_ring': stalk_color_below_rings,
                  'veil_type': veil_types,
                  'veil_color': veil_colors,
                  'ring_number': ring_numbers,
                  'ring_type': ring_types,
                  'spore_print_color': spore_print_colors,
                  'population': populations,
                  'habitat': habitats},
                inplace=True)
```



0.2 Data Visualization



```
[10]: data['classes'] = np.where(data['classes'] == 'poisonous', 1, 0)
```

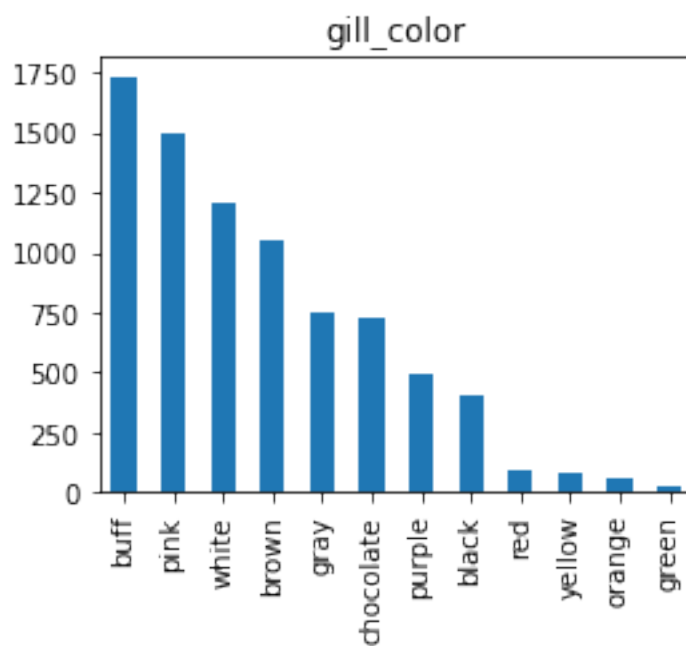

```
[11]: for col in data.columns:
        plt.figure(figsize = (4,3))
        data[col].value_counts().plot(kind='bar')
        plt.title(col)
        plt.show()
```






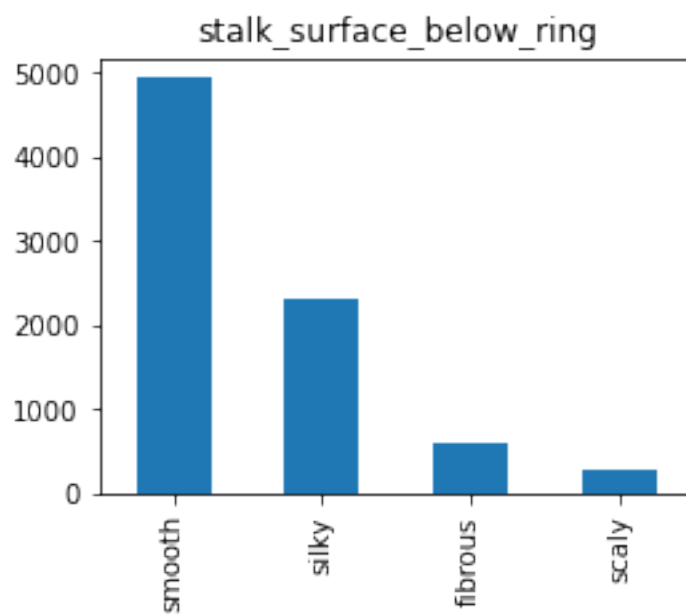




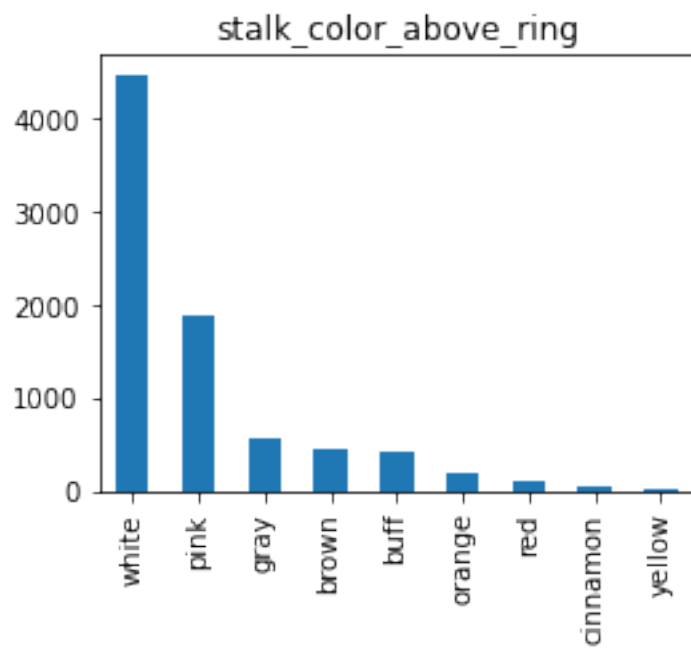



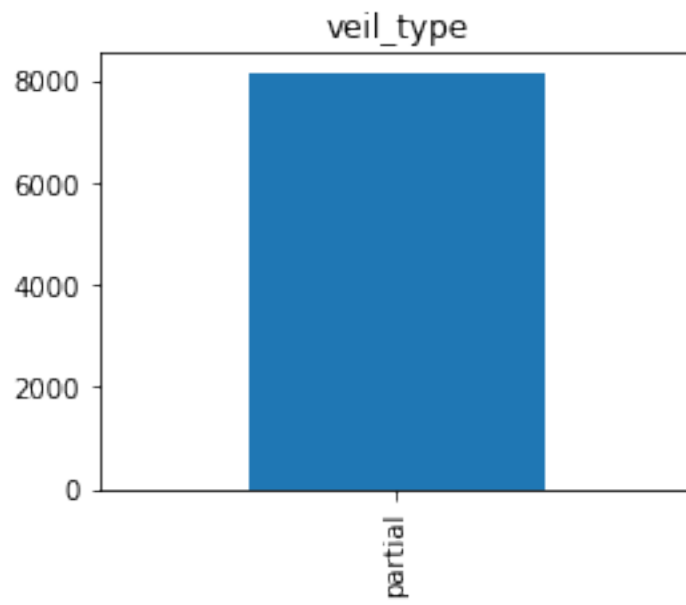


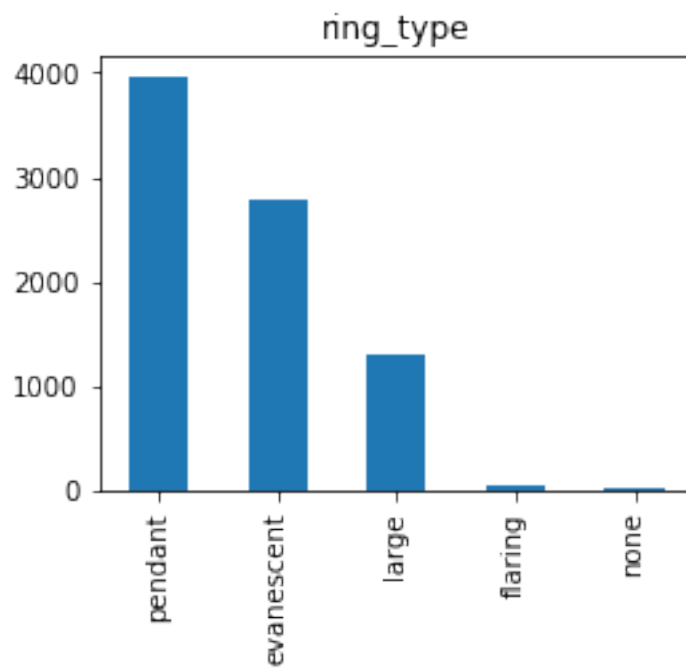
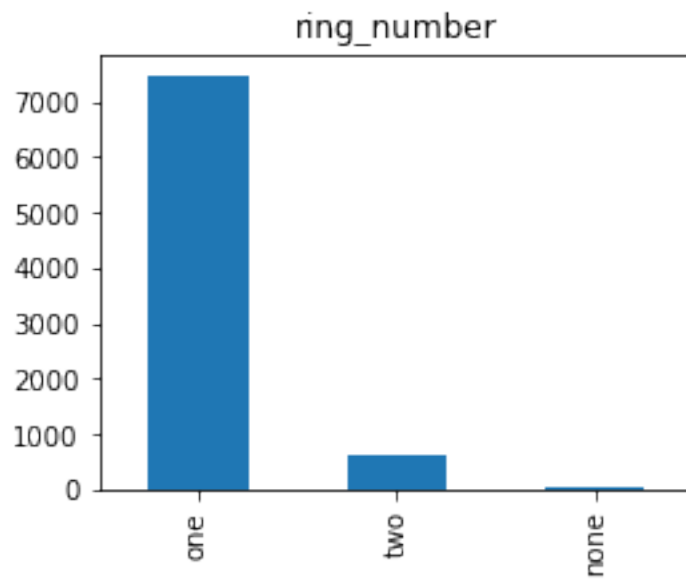



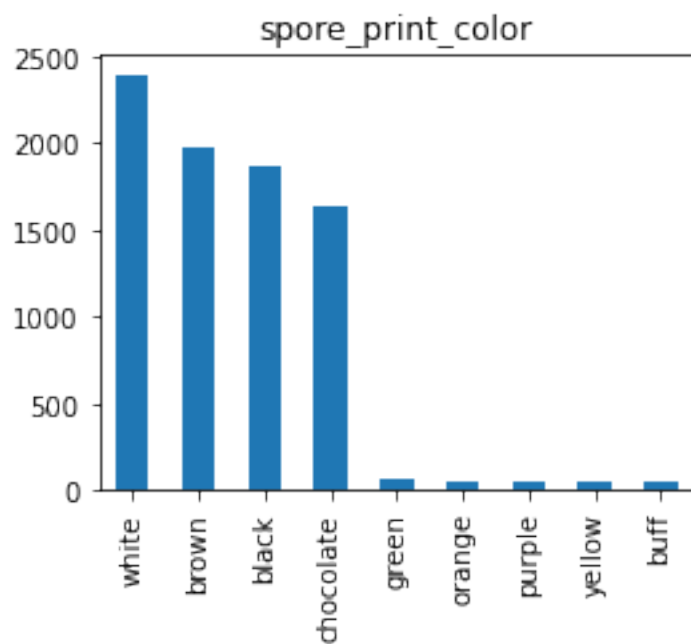


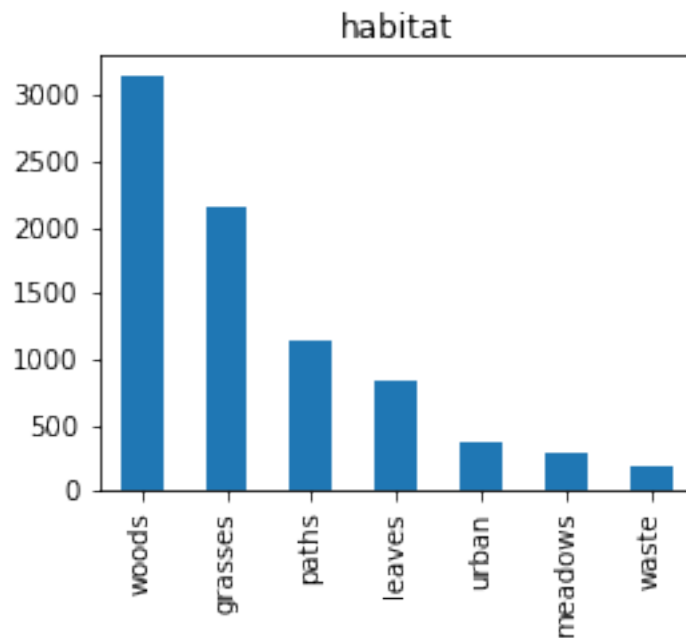













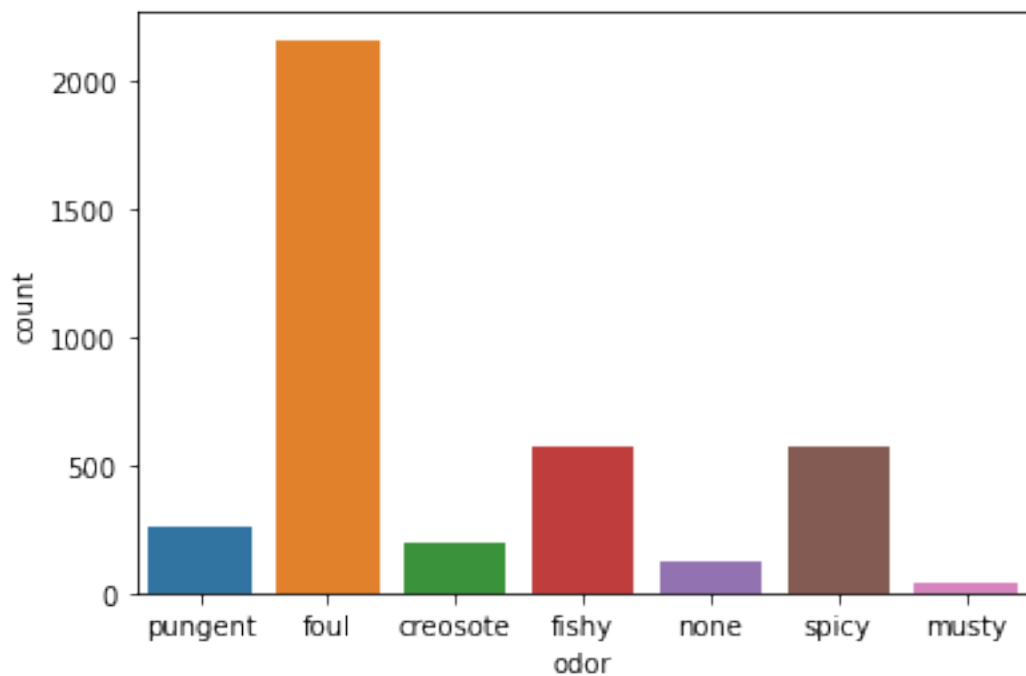
```
[12]: plt.figure(figsize = (8,6))  
sns.countplot(data = data, x = 'odor', hue = 'classes')
```

```
[12]: <AxesSubplot:xlabel='odor', ylabel='count'>
```



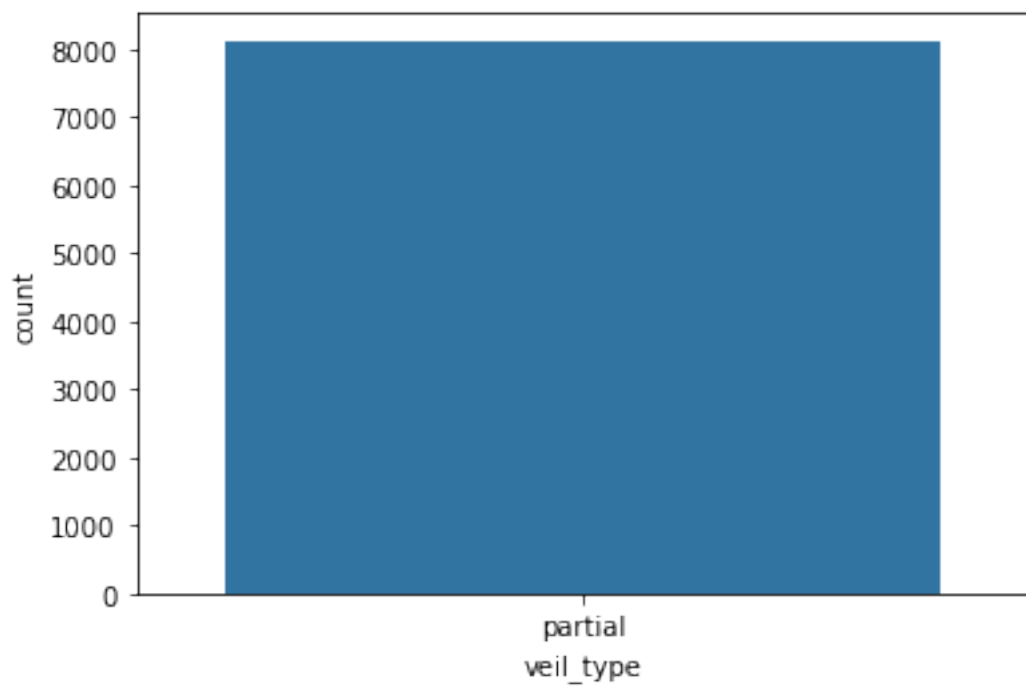
```
[13]: sns.countplot(data = data[data['classes'] == 1], x = 'odor')
```

```
[13]: <AxesSubplot:xlabel='odor', ylabel='count'>
```



```
[14]: sns.countplot(data = data, x = 'veil_type')
```


```
[14]: <AxesSubplot:xlabel='veil_type', ylabel='count'>
```



```
[15]: poison_df = data[data['classes'] == 1]
      edible_df = data[data['classes'] == 0]
```


```
[16]: plt.figure(figsize = (8,6))
      sns.countplot(data=poison_df, x='gill_color')
```

```
[16]: <AxesSubplot:xlabel='gill_color', ylabel='count'>
```



```
[17]: plt.figure(figsize = (8,6))
      sns.countplot(data=edible_df, x='gill_color')
```

```
[17]: <AxesSubplot:xlabel='gill_color', ylabel='count'>
```



0.3 Confusion Matrix Functions

```
[18]: #Function to compute confusion matrix and prediction accuracy on test/train
      ↪data -- Decision Tree
def confusion_matrix_data(data,actual_values,model,cutoff=0.5):
    #Predict the values using the Logit model
    pred_values = model.predict_proba(data)[: ,1]
    # Specify the bins
    bins=np.array([0,cutoff,1])
    #Confusion matrix
    cm = np.histogram2d(actual_values, pred_values, bins=bins)[0]
    cm_df = pd.DataFrame(cm)
    cm_df.columns = ['Predicted 0','Predicted 1']
    cm_df = cm_df.rename(index={0: 'Actual 0',1:'Actual 1'})
    # Calculate the accuracy
    accuracy = 100*(cm[0,0]+cm[1,1])/cm.sum()
    fpr = 100*(cm[0,1])/(cm[0,1]+cm[0,0])
    fnr = 100*(cm[1,0])/(cm[1,0]+cm[1,1])
    recall = 100*(cm[1,1])/(cm[1,0]+cm[1,1])
    print("Accuracy = ", accuracy)
    print("Recall = ", recall)
```



```

print("FPR = ", fpr)
print("FNR = ", fnr)
print("Confusion matrix = \n", cm_df)
return (" ")

```

```

[19]: #Function to compute confusion matrix and prediction accuracy on test/train
      ↪data -- Decision Tree
def confusion_matrix_data_logit(data,actual_values,model,cutoff=0.5):
    #Predict the values using the Logit model
    pred_values = model.predict(data)
    # Specify the bins
    bins=np.array([0,cutoff,1])
    #Confusion matrix
    cm = np.histogram2d(actual_values, pred_values, bins=bins)[0]
    cm_df = pd.DataFrame(cm)
    cm_df.columns = ['Predicted 0','Predicted 1']
    cm_df = cm_df.rename(index={0: 'Actual 0',1:'Actual 1'})
    # Calculate the accuracy
    accuracy = 100*(cm[0,0]+cm[1,1])/cm.sum()
    fnr = 100*(cm[1,0])/(cm[1,0]+cm[1,1])
    print("Accuracy = ", accuracy)
    print("FNR = ", fnr)
    print("Confusion matrix = \n", cm_df)
    return (" ")

```

```

[20]: from sklearn.model_selection import train_test_split
      train, test = train_test_split(data, test_size=0.33, random_state=1)

```

```

[21]: print(train.shape)
      print(test.shape)

```

```

(5443, 23)

```

```

(2681, 23)

```

0.4 Linear model – LOGISTIC REGRESSION

```

[22]: train.columns.to_list()

```

```

[22]: ['classes',
      'cap_shape',
      'cap_surface',
      'cap_color',
      'bruises',
      'odor',
      'gill_attachment',
      'gill_spacing',
      'gill_size',
      'gill_color',

```

```
'stalk_shape',
'stalk_root',
'stalk_surface_above_ring',
'stalk_surface_below_ring',
'stalk_color_above_ring',
'stalk_color_below_ring',
'veil_type',
'veil_color',
'ring_number',
'ring_type',
'spore_print_color',
'population',
'habitat']
```

```
[23]: train_m1 = train[['classes',
    'cap_shape',
    'cap_surface',
    'cap_color',
    'bruises',
    'gill_attachment',
    'gill_spacing',
    'gill_size',
    'gill_color',
    'stalk_shape',
    'stalk_root',
    'stalk_surface_above_ring',
    'stalk_surface_below_ring',
    'stalk_color_above_ring',
    'stalk_color_below_ring',
    'veil_color',
    'ring_number',
    'ring_type',
    'spore_print_color',
    'population',
    'habitat']]
```

```
test_m1 = test[['classes',
    'cap_shape',
    'cap_surface',
    'cap_color',
    'bruises',
    'gill_attachment',
    'gill_spacing',
    'gill_size',
    'gill_color',
    'stalk_shape',
    'stalk_root',
```

```
'stalk_surface_above_ring',
'stalk_surface_below_ring',
'stalk_color_above_ring',
'stalk_color_below_ring',
'veil_color',
'ring_number',
'ring_type',
'spore_print_color',
'population',
'habitat']]
```

```
[24]: X=data.drop(['classes', 'odor', 'veil_type'],axis=1) #Predictors
      y=data['classes'] #Response
```

```
[25]: from sklearn.preprocessing import LabelEncoder
      Encoder_X = LabelEncoder()
      for col in X.columns:
          X[col] = Encoder_X.fit_transform(X[col])
      Encoder_y=LabelEncoder()
      y = Encoder_y.fit_transform(y)
```

```
[26]: X=pd.get_dummies(X,columns=X.columns,drop_first=True)
```

```
[27]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      ↪random_state=42)
```

```
[28]: from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()

      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

```
[29]: from sklearn.linear_model import LogisticRegression
      classifier = LogisticRegression()

      classifier.fit(X_train,y_train)
```

```
[29]: LogisticRegression()
```

```
[30]: from sklearn.metrics import confusion_matrix,classification_report
      from sklearn.model_selection import cross_val_predict, cross_val_score
      from sklearn.metrics import
      ↪confusion_matrix,classification_report,accuracy_score
```

```
[31]: def print_score(classifier,X_train,y_train,X_test,y_test,train=True):
      if train == True:
```

```

    print("Training results:\n")
    print('Accuracy Score: {0:.4f}\n'.
    ↪format(accuracy_score(y_train,classifier.predict(X_train))))
    print('Classification Report:\n{}\n'.
    ↪format(classification_report(y_train,classifier.predict(X_train))))
    print('Confusion Matrix:\n{}\n'.
    ↪format(confusion_matrix(y_train,classifier.predict(X_train))))
    res = cross_val_score(classifier, X_train, y_train, cv=10, n_jobs=-1,
    ↪scoring='accuracy')
    print('Average Accuracy:\t{0:.4f}\n'.format(res.mean()))
    print('Standard Deviation:\t{0:.4f}'.format(res.std()))
    elif train == False:
        print("Test results:\n")
        print('Accuracy Score: {0:.4f}\n'.
        ↪format(accuracy_score(y_test,classifier.predict(X_test))))
        print('Classification Report:\n{}\n'.
        ↪format(classification_report(y_test,classifier.predict(X_test))))
        print('Confusion Matrix:\n{}\n'.
        ↪format(confusion_matrix(y_test,classifier.predict(X_test))))

```

```
[32]: print_score(classifier,X_train,y_train,X_test,y_test,train=True)
```

Training results:

Accuracy Score: 1.0000

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2951
1	1.00	1.00	1.00	2735
accuracy			1.00	5686
macro avg	1.00	1.00	1.00	5686
weighted avg	1.00	1.00	1.00	5686

Confusion Matrix:

```
[[2951   0]
 [   0 2735]]
```

Average Accuracy: 0.9998

Standard Deviation: 0.0005

0.5 Non-linear model – DECISION TREE

```
[33]: train_m2 = pd.get_dummies(train)
test_m2 = pd.get_dummies(test)
```

```
[34]: X2 = train_m2.drop(columns = 'classes')
      X2test = test_m2.drop(columns = 'classes')
      y2 = train_m2['classes']
      y2test = test_m2['classes']
```

```
[35]: model2 = DecisionTreeClassifier(random_state=1, max_depth=3)
      model2.fit(X2, y2)
```

```
[35]: DecisionTreeClassifier(max_depth=3, random_state=1)
```

```
[36]: confusion_matrix_data(X2, train_m2.classes, model2, cutoff=0.5)
```

```
Accuracy = 98.40161675546574
Recall = 99.88461538461539
FPR = 2.9546253957087583
FNR = 0.11538461538461539
Confusion matrix =
```

	Predicted 0	Predicted 1
Actual 0	2759.0	84.0
Actual 1	3.0	2597.0

[36] : ' '

```
[37]: X2.columns[pd.Series(model2.feature_importances_ > 0)]
```

```
[37]: Index(['odor_none', 'stalk_root_club', 'stalk_surface_below_ring_scaly',
            'stalk_surface_below_ring_smooth', 'spore_print_color_green'],
           dtype='object')
```

```
[38]: model2.feature_importances_
```

```
[38]: array([0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.6656145, 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.1827068, 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.09808967, 0.      , 0.01753364, 0.      ,
```

0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.03605539,	
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.])			

0.6 Data Cleaning for further models

```
[39]: train.columns.to_list()
```

```
[39]: ['classes',
       'cap_shape',
       'cap_surface',
       'cap_color',
       'bruises',
       'odor',
       'gill_attachment',
       'gill_spacing',
       'gill_size',
       'gill_color',
       'stalk_shape',
       'stalk_root',
       'stalk_surface_above_ring',
       'stalk_surface_below_ring',
       'stalk_color_above_ring',
       'stalk_color_below_ring',
       'veil_type',
       'veil_color',
       'ring_number',
       'ring_type',
       'spore_print_color',
       'population',
       'habitat']
```

```
[40]: train.gill_spacing.value_counts()
```

```
[40]: close      4544
      crowded     899
      Name: gill_spacing, dtype: int64
```

```
[41]: train.gill_size.value_counts()
```

```
[41]: broad      3771
      narrow    1672
      Name: gill_size, dtype: int64
```

```
[42]: train.stalk_root.value_counts()
```

```
[42]: bulbous      2555
      NA          1653
      equal       750
      club        362
      rooted      123
      Name: stalk_root, dtype: int64
```

```
[43]: train.stalk_surface_above_ring.value_counts()
```

```
[43]: smooth      3477
      silky      1584
      fibrous     366
      scaly       16
      Name: stalk_surface_above_ring, dtype: int64
```

```
[44]: train.population.value_counts()
```

```
[44]: several      2724
      solitary    1143
      scattered    822
      numerous     278
      abundant     257
      clustered    219
      Name: population, dtype: int64
```

0.7 RANDOM FOREST

```
[45]: train_m3 = pd.get_dummies(train[['classes', 'cap_shape', 'cap_surface',
    ↪ 'cap_color', 'habitat', 'stalk_shape']])
      test_m3 = pd.get_dummies(test[['classes', 'cap_shape', 'cap_surface',
    ↪ 'cap_color', 'habitat', 'stalk_shape']])
```

```
[46]: X = train_m3.drop(columns = 'classes')
      Xtest = test_m3.drop(columns = 'classes')
      y = train_m3['classes']
      ytest = test_m3['classes']
```

```
[47]: params = {'n_estimators': [500],
               'max_features': range(1,6),
               }

      param_list = list(it.product(*(params[Name] for Name in list(params.keys()))))
```

```

recall = [0]*len(param_list)

i=0
for pr in param_list:
    model = RandomForestClassifier(random_state=1,
                                   oob_score=True,
                                   verbose=False,
                                   n_estimators = pr[0],
                                   max_features=pr[1],
                                   n_jobs=-1).fit(X,y)

    oob_pred = model.oob_decision_function_[:,1]
    bins=np.array([0,0.5,1])
    cm = np.histogram2d(y, oob_pred, bins=bins)[0]
    recall[i] = 100*(cm[1,1])/(cm[1,0]+cm[1,1])
    i=i+1

end_time = time.time()

print("max recall = ", np.max(recall))
print("params= ", param_list[np.argmax(recall)])

```

```

max recall = 90.5
params= (500, 1)

```

```
[48]: model3 = RandomForestClassifier(random_state=1, n_jobs=-1, max_features=1,
↪n_estimators=100).fit(X, y)
```

```
[49]: confusion_matrix_data(X, y, model3, cutoff=0.5)
```

```

Accuracy = 93.7167003490722
Recall = 90.84615384615384
FPR = 3.658107632782272
FNR = 9.153846153846153
Confusion matrix =

```

	Predicted 0	Predicted 1
Actual 0	2739.0	104.0
Actual 1	238.0	2362.0

```
[49]: ' '
```

```
[50]: confusion_matrix_data(Xtest, ytest, model3, cutoff=0.5)
```

```

Accuracy = 92.83849309958971
Recall = 91.48936170212765
FPR = 5.86080586080586
FNR = 8.51063829787234
Confusion matrix =

```

	Predicted 0	Predicted 1
Actual 0	1285.0	80.0

Actual 1	112.0	1204.0
----------	-------	--------

[50]: ' '

0.8 Tuned Decision Tree – Rule of Five

```
[51]: param_grid = {
        'max_depth': range(1,5),
        'max_leaf_nodes': range(1,30),
        'max_features': range(1,5),
    }

    skf = StratifiedKFold(n_splits=5)

    grid_search = GridSearchCV(DecisionTreeClassifier(random_state=1),
                                param_grid,
                                scoring=['precision', 'recall', 'accuracy'],
                                refit="recall",
                                cv=skf,
                                n_jobs=-1,
                                verbose = True).fit(X,y)

    print('Best params for recall')
    print(grid_search.best_params_)
```

Fitting 5 folds for each of 464 candidates, totalling 2320 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.

[Parallel(n_jobs=-1)]: Done 26 tasks | elapsed: 0.7s

[Parallel(n_jobs=-1)]: Done 762 tasks | elapsed: 2.2s

Best params for recall

{'max_depth': 4, 'max_features': 4, 'max_leaf_nodes': 5}

[Parallel(n_jobs=-1)]: Done 2320 out of 2320 | elapsed: 4.8s finished

```
[52]: model4 = DecisionTreeClassifier(random_state=1, max_depth = 4, max_features = 4,
    ↪max_leaf_nodes=5)

    model4.fit(X,y)
    print(confusion_matrix_data(X,y,model4))
    print(confusion_matrix_data(Xtest,ytest,model4,cutoff=0.4))
```

Accuracy = 59.02994672055851

Recall = 88.53846153846153

FPR = 67.95638410130144

FNR = 11.461538461538462

Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	911.0	1932.0

```
Actual 1          298.0          2302.0
```

```
Accuracy = 59.716523685192094
```

```
Recall = 95.82066869300913
```

```
FPR = 75.0915750915751
```

```
FNR = 4.179331306990881
```

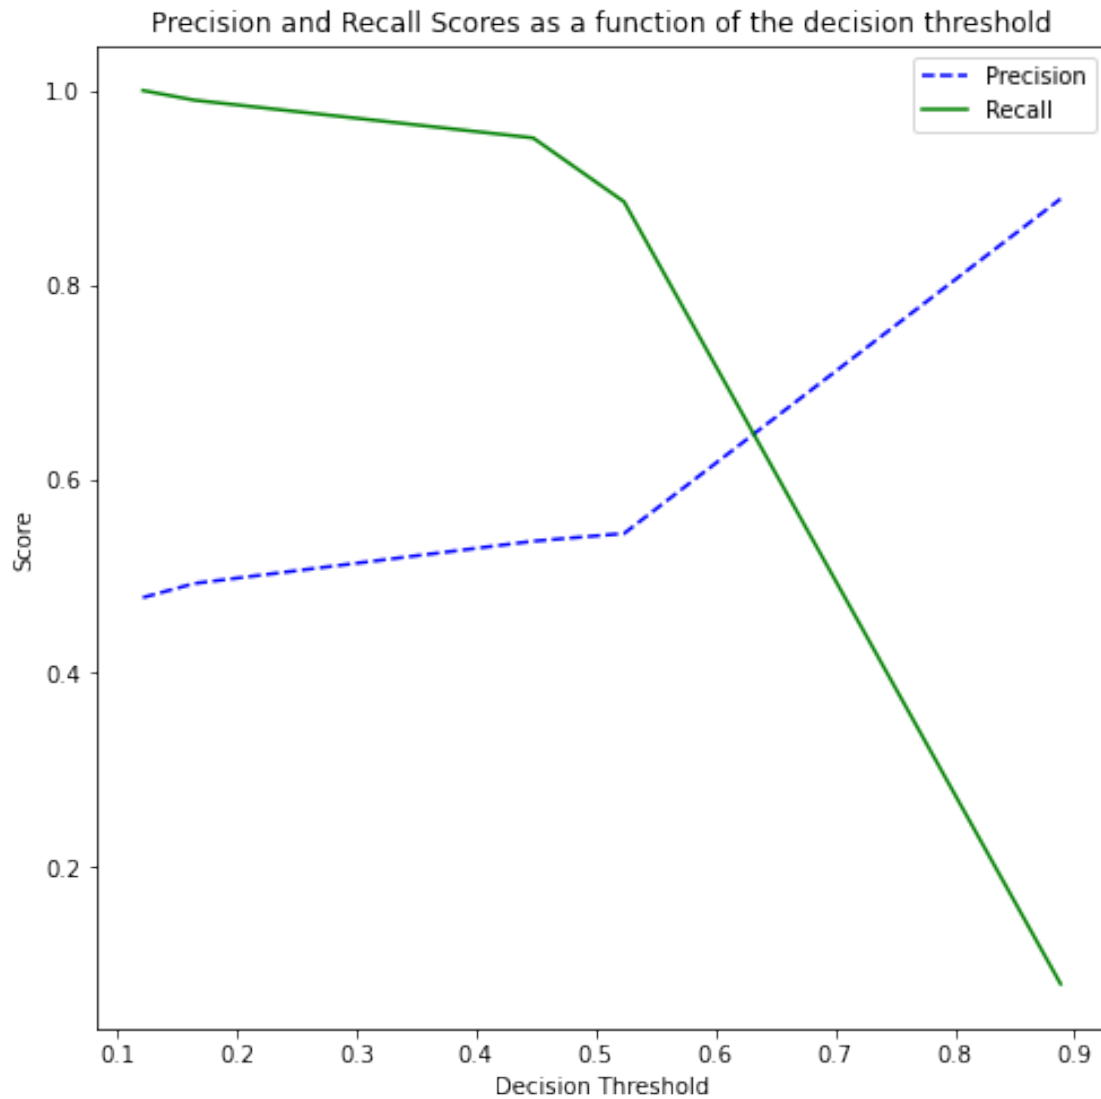
```
Confusion matrix =
```

	Predicted 0	Predicted 1
Actual 0	340.0	1025.0
Actual 1	55.0	1261.0

```
[53]: ypred = model4.predict_proba(X)[: , 1]
      p, r, thresholds = precision_recall_curve(y, ypred)
```

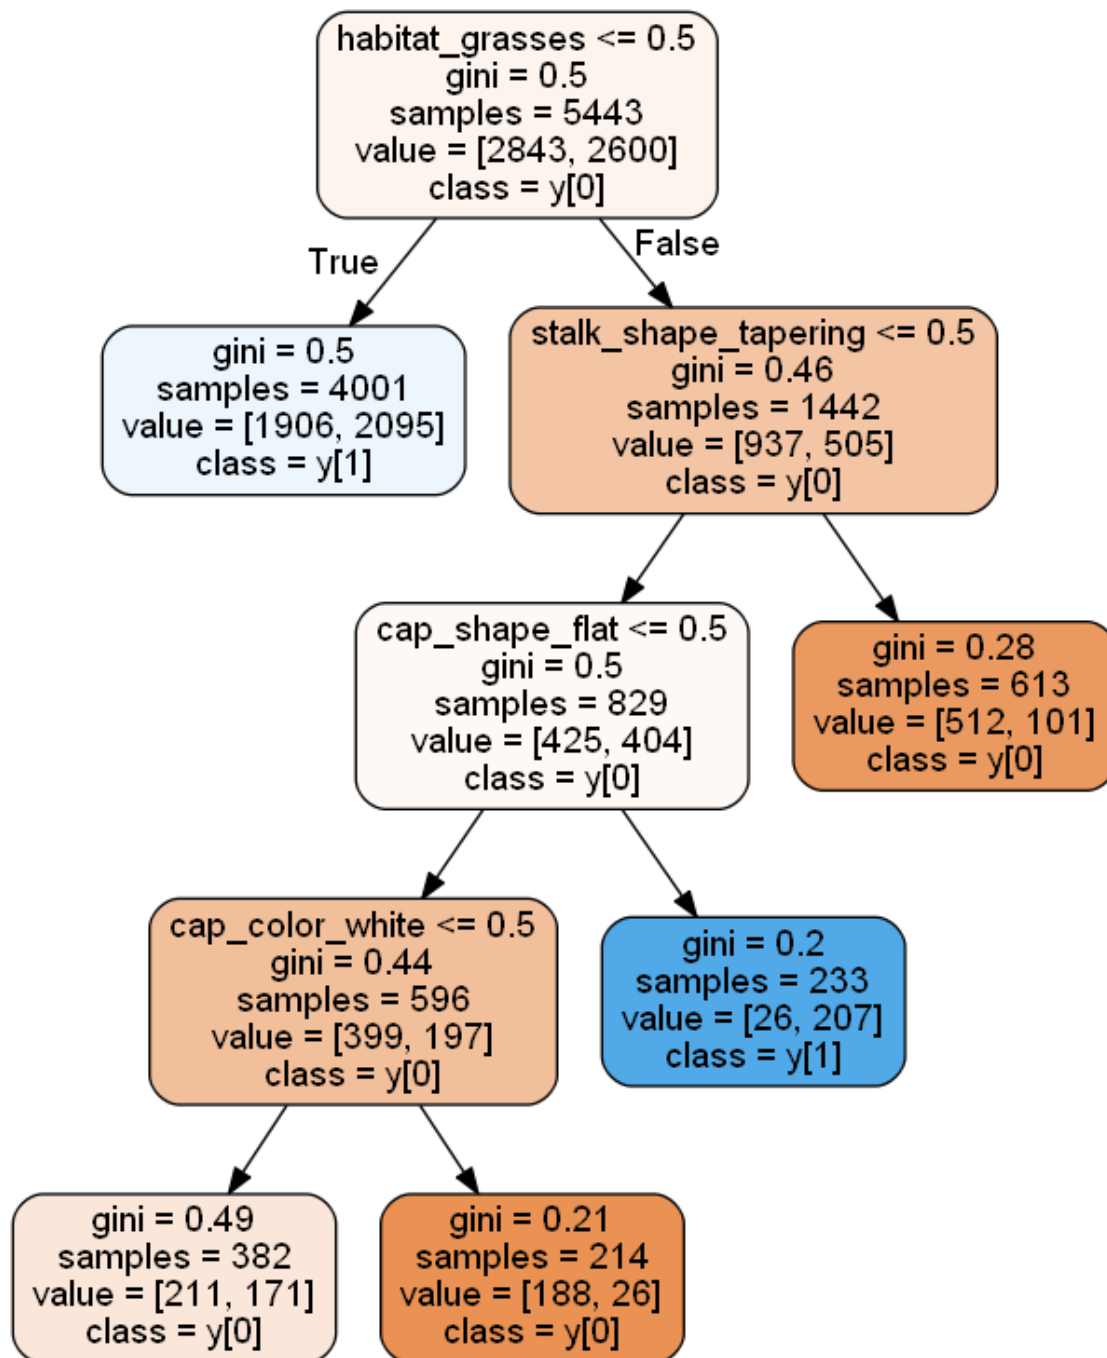
```
[54]: def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
      plt.figure(figsize=(8, 8))
      plt.title("Precision and Recall Scores as a function of the decision_
      ↪threshold")
      plt.plot(thresholds, precisions[:-1], "b--", label="Precision")
      plt.plot(thresholds, recalls[:-1], "g-", label="Recall")
      plt.ylabel("Score")
      plt.xlabel("Decision Threshold")
      plt.legend(loc='best')

      plot_precision_recall_vs_threshold(p, r, thresholds)
```



```
[55]: #Visualizing the regression tree
dot_data = StringIO()
export_graphviz(model4, out_file=dot_data,
                filled=True, rounded=True, class_names=True,
                feature_names =X.columns,precision=2)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png('car_price_tree.png')
Image(graph.create_png())
```

[55]:



0.8.1 Finding most accurate prediction subset

```
[56]: #The following function finds indices of the leaf nodes, given the decision
      ↪ tree model.
      #The indices of nodes in a decision tree starts at the top from 0.
```

```

#The child nodes obtained from the first split have indices 1 (left) and 2
    ↪ (right),
#the child nodes obtained from the second split have indices 3 (left) and 4
    ↪ (right), and so on.
def leaf_nodes_indices(model):
    children_left = model.tree_.children_left
    children_right = model.tree_.children_right
    leaf_nodes = []
    for i in range(model.tree_.node_count):
        if children_left[i] == children_right[i]:
            leaf_nodes.append(i)
    return leaf_nodes

#The following function finds the *mean squared error* of nodes with indices
    ↪ *node_indices*
def mse(model,node_indices):
    return model.tree_.impurity[node_indices]

#The following function gives the decision rules for a node with index as
    ↪ *node_index*
def decision_rules(model,node_index):
    child_node = node_index
    node_list=[]
    children_left = model.tree_.children_left
    children_right = model.tree_.children_right
    features=model.tree_.feature
    fnames = X.columns
    threshold = model.tree_.threshold
    p=1
    while p>0:
        if node_index%2>0:
            p = np.where(children_left==node_index)[0][0]
        else:
            p = np.where(children_right==node_index)[0][0]
        node_list.append(p)
        node_index=p
    node_list.reverse()
    node_list.append(child_node)
    cc=1
    for n in node_list[0:(len(node_list)-1)]:
        cnode = node_list[cc]
        if cnode%2==0:
            ineq_sign = ">"
        else:
            ineq_sign = "<="
        print("Split "+ str(cc)+" :
    ↪ "+fnames[features][n]+ineq_sign+str(threshold[n]))

```

```

        cc=cc+1
        node_list=[]
        return ""

```

```

[57]: nodes = leaf_nodes_indices(model4)
      m = nodes[0]
      for i in nodes:
          if mse(model4, i)< m:
              m=i
      print("leaf=",m, "mse=", mse(model4, m))

```

```

leaf= 8 mse= 0.2134684251899729

```

```

[59]: print(decision_rules(model4, 8))

```

```

Split 1:habitat_grasses>0.5
Split 2:stalk_shape_tapering<=0.5
Split 3:cap_shape_flat<=0.5
Split 4:cap_color_white>0.5

```

```

[60]: test_filtered = test_m3[(test_m3.habitat_grasses<=0.5) &
                              (test_m3.stalk_shape_tapering<=0.5) &
                              (test_m3.cap_color_brown<=0.5) &
                              (test_m3.cap_surface_smooth>0.5) &
                              (test_m3.cap_color_pink<=0.5) &
                              (test_m3.habitat_waste<=0.5) &
                              (test_m3.cap_color_white>0.5) &
                              (test_m3.habitat_meadows>0.5)]

```

```

[61]: Xtest_filt = test_filtered.drop(columns = 'classes')
      ytest_filt = test_filtered['classes']

```

```

[62]: print(confusion_matrix_data(Xtest_filt,ytest_filt,model4))

```

```

Accuracy = 8.695652173913043
Recall = 100.0
FPR = 100.0
FNR = 0.0
Confusion matrix =

```

	Predicted 0	Predicted 1
Actual 0	0.0	21.0
Actual 1	0.0	2.0

0.9 Tuned Decision Tree (more predictors) - Max Depth 3

```
[63]: train_m4 = pd.get_dummies(train.drop(columns = ['odor', 'veil_type']))
      test_m4 = pd.get_dummies(test.drop(columns = ['odor', 'veil_type']))
```

```
[64]: X = train_m4.drop(columns = 'classes')
      Xtest = test_m4.drop(columns = 'classes')
      y = train_m4['classes']
      ytest = test_m4['classes']
```

```
[65]: param_grid = {
      'max_leaf_nodes': range(10,40),
      'max_features': range(1,108),
      'max_depth': range(1,9)
      }

      skf = StratifiedKFold(n_splits=5)

      grid_search = GridSearchCV(DecisionTreeClassifier(random_state=1),
                                param_grid,
                                scoring=['precision', 'recall', 'accuracy'],
                                refit="recall",
                                cv=skf,
                                n_jobs=-1,
                                verbose = True).fit(X,y)

      print('Best params for recall')
      print(grid_search.best_params_)
```

Fitting 5 folds for each of 25680 candidates, totalling 128400 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 28 tasks      | elapsed:    0.0s
[Parallel(n_jobs=-1)]: Done 1000 tasks    | elapsed:    2.3s
[Parallel(n_jobs=-1)]: Done 3000 tasks    | elapsed:    6.7s
[Parallel(n_jobs=-1)]: Done 5800 tasks    | elapsed:   13.3s
[Parallel(n_jobs=-1)]: Done 9400 tasks    | elapsed:   22.3s
[Parallel(n_jobs=-1)]: Done 13800 tasks   | elapsed:   34.3s
[Parallel(n_jobs=-1)]: Done 19000 tasks   | elapsed:   48.3s
[Parallel(n_jobs=-1)]: Done 25000 tasks   | elapsed:   1.1min
[Parallel(n_jobs=-1)]: Done 31800 tasks   | elapsed:   1.4min
[Parallel(n_jobs=-1)]: Done 39400 tasks   | elapsed:   1.8min
[Parallel(n_jobs=-1)]: Done 47800 tasks   | elapsed:   2.3min
[Parallel(n_jobs=-1)]: Done 57000 tasks   | elapsed:   2.7min
[Parallel(n_jobs=-1)]: Done 67000 tasks   | elapsed:   3.3min
[Parallel(n_jobs=-1)]: Done 77800 tasks   | elapsed:   4.0min
[Parallel(n_jobs=-1)]: Done 89400 tasks   | elapsed:   4.7min
[Parallel(n_jobs=-1)]: Done 101800 tasks  | elapsed:   5.4min
[Parallel(n_jobs=-1)]: Done 115000 tasks  | elapsed:   6.2min
```

[Parallel(n_jobs=-1)]: Done 128400 out of 128400 | elapsed: 7.1min finished

Best params for recall

{'max_depth': 5, 'max_features': 2, 'max_leaf_nodes': 10}

```
[66]: model5 = DecisionTreeClassifier(random_state=1, max_features = 2,
    ↪max_leaf_nodes=10, max_depth = 5)
```

```
model5.fit(X,y)
print(confusion_matrix_data(X,y,model5))
print(confusion_matrix_data(Xtest,ytest,model5))
```

Accuracy = 75.27099026272276

Recall = 100.0

FPR = 47.34435455504749

FNR = 0.0

Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	1497.0	1346.0
Actual 1	0.0	2600.0

Accuracy = 75.60611712047744

Recall = 100.0

FPR = 47.91208791208791

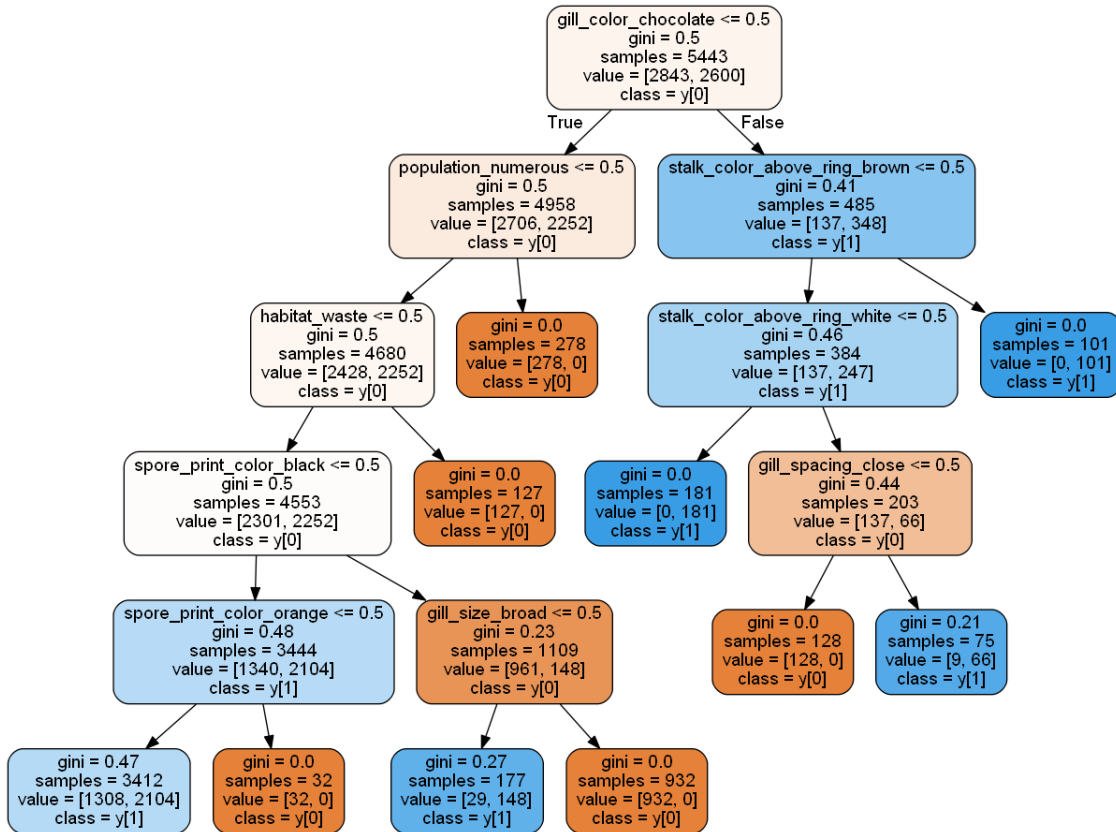
FNR = 0.0

Confusion matrix =

	Predicted 0	Predicted 1
Actual 0	711.0	654.0
Actual 1	0.0	1316.0

```
[67]: #Visualizing the regression tree
dot_data = StringIO()
export_graphviz(model5, out_file=dot_data,
                filled=True, rounded=True, class_names = True,
                feature_names =X.columns,precision=2)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
#graph.write_png('car_price_tree.png')
Image(graph.create_png())
```

[67]:



```
[68]: X.columns[pd.Series(model5.feature_importances_ > 0)]
```

```
[68]: Index(['gill_spacing_close', 'gill_size_broad', 'gill_color_chocolate',
           'stalk_color_above_ring_brown', 'stalk_color_above_ring_white',
           'spore_print_color_black', 'spore_print_color_orange',
           'population_numerous', 'habitat_waste'],
          dtype='object')
```

```
[69]: train_m5 = pd.get_dummies(train[['classes', 'gill_spacing', 'gill_size',
    ↪ 'gill_color', 'stalk_color_above_ring',
    ↪ 'spore_print_color', 'population', 'habitat']])
test_m5 = pd.get_dummies(test[['classes', 'gill_spacing', 'gill_size',
    ↪ 'gill_color', 'stalk_color_above_ring',
    ↪ 'spore_print_color', 'population', 'habitat']])
```

```
[70]: X = train_m5.drop(columns = 'classes')
Xtest = test_m5.drop(columns = 'classes')
y = train_m5['classes']
ytest = test_m5['classes']
```

```
[71]: param_grid = {
        'max_leaf_nodes': range(10,40),
        'max_features': range(1,48),
        'max_depth': range(1,9)
    }

    skf = StratifiedKFold(n_splits=5)

    grid_search = GridSearchCV(DecisionTreeClassifier(random_state=1),
                                param_grid,
                                scoring=['precision','recall','accuracy'],
                                refit="recall",
                                cv=skf,
                                n_jobs=-1,
                                verbose = True).fit(X,y)

    print('Best params for recall')
    print(grid_search.best_params_)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 28 tasks      | elapsed:    0.0s
```

Fitting 5 folds for each of 11280 candidates, totalling 56400 fits

```
[Parallel(n_jobs=-1)]: Done 1000 tasks      | elapsed:    2.1s
[Parallel(n_jobs=-1)]: Done 3000 tasks      | elapsed:    6.3s
[Parallel(n_jobs=-1)]: Done 5800 tasks      | elapsed:   12.2s
[Parallel(n_jobs=-1)]: Done 9400 tasks      | elapsed:   19.7s
[Parallel(n_jobs=-1)]: Done 13800 tasks     | elapsed:   29.1s
[Parallel(n_jobs=-1)]: Done 19000 tasks     | elapsed:   41.4s
[Parallel(n_jobs=-1)]: Done 25000 tasks     | elapsed:   55.0s
[Parallel(n_jobs=-1)]: Done 31800 tasks     | elapsed:  1.2min
[Parallel(n_jobs=-1)]: Done 39400 tasks     | elapsed:  1.5min
[Parallel(n_jobs=-1)]: Done 47800 tasks     | elapsed:  1.8min
[Parallel(n_jobs=-1)]: Done 56400 out of 56400 | elapsed:  2.2min finished
```

Best params for recall

```
{'max_depth': 6, 'max_features': 25, 'max_leaf_nodes': 15}
```

```
[72]: model6 = DecisionTreeClassifier(random_state=1, max_features = 25,
    ↪max_leaf_nodes=15, max_depth = 6)

    model6.fit(X,y)
    print(confusion_matrix_data(X,y,model6))
    print(confusion_matrix_data(Xtest,ytest,model6))
```

Accuracy = 99.11813338232592

Recall = 99.46153846153847

FPR = 1.1959198030249736

FNR = 0.5384615384615384

```
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      2809.0      34.0
Actual 1      14.0      2586.0
```

Accuracy = 99.10481163744872

Recall = 99.24012158054711

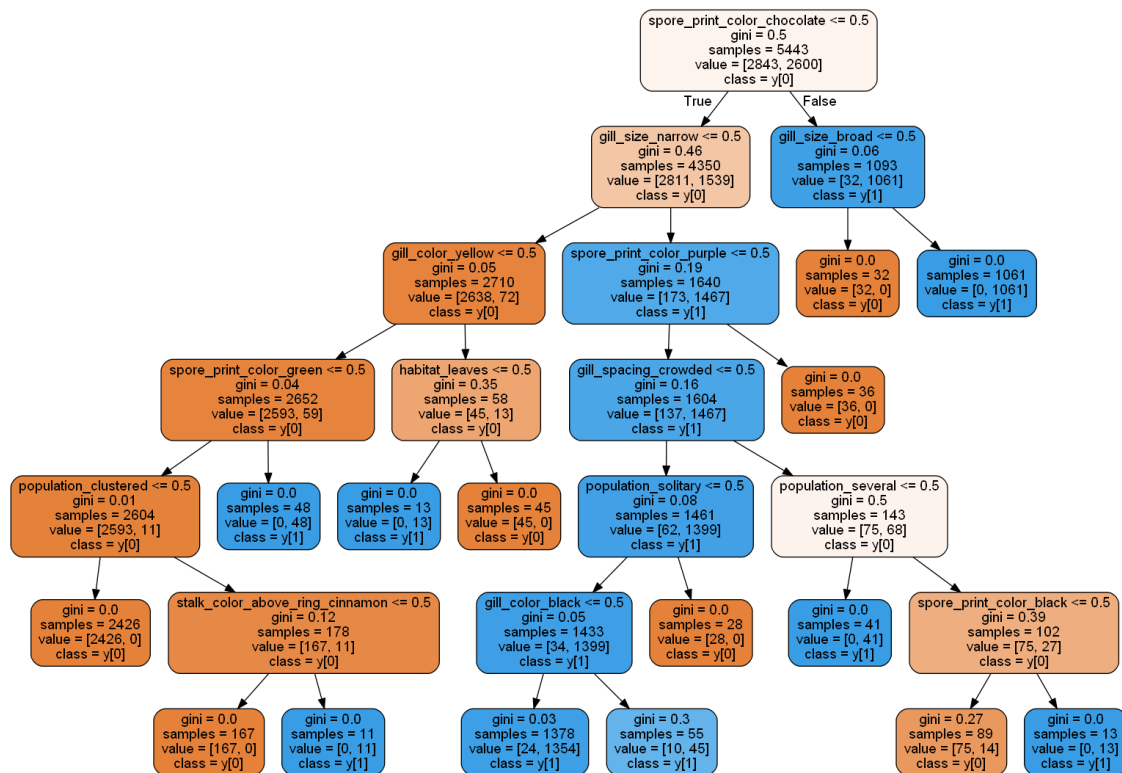
FPR = 1.0256410256410255

FNR = 0.7598784194528876

```
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      1351.0      14.0
Actual 1      10.0      1306.0
```

```
[73]: #Visualizing the decision tree
dot_data = StringIO()
export_graphviz(model6, out_file=dot_data,
                filled=True, rounded=True, class_names = True,
                feature_names =X.columns,precision=2)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
#graph.write_png('car_price_tree.png')
Image(graph.create_png())
```

[73]:



0.10 ENSEMBLE MODEL

0.10.1 Decision Tree

```
[74]: m1 = model6
```

0.10.2 Random Forest

```
[75]: params = {'n_estimators': [500],
               'max_features': range(1,48),
               }

param_list=list(it.product(*(params[Name] for Name in list(params.keys()))))
recall = [0]*len(param_list)

i=0
for pr in param_list:
    model =
    ↪RandomForestClassifier(random_state=1,oob_score=True,verbose=False,n_estimators=
    ↪ pr[0],
                                max_features=pr[1], n_jobs=-1).fit(X,y)
    oob_pred = model.oob_decision_function_[:,1]
    bins=np.array([0,0.5,1])
    cm = np.histogram2d(y, oob_pred, bins=bins)[0]
    recall[i] = 100*(cm[1,1])/(cm[1,0]+cm[1,1])
    i=i+1

print("max recall = ", np.max(recall))
print("params= ", param_list[np.argmax(recall)])
```

```
max recall = 99.3076923076923
params= (500, 1)
```

```
[76]: m2 =
    ↪RandomForestClassifier(random_state=1,n_jobs=-1,max_features=1,n_estimators=500).
    ↪fit(X, y)
```

0.10.3 XGBoost

```
[77]: param_grid = {'n_estimators': [25,100,500],
                   'max_depth': [6,7,8],
                   'learning_rate': [0.01,0.1,0.2],
                   'gamma': [0.1,0.25,0.5],
                   'reg_lambda': [0,0.01,0.001],
                   'scale_pos_weight': [1.25,1.5,1.75]
                   }

cv = StratifiedKFold(n_splits=5,shuffle=True,random_state=1)
```

```

optimal_params = GridSearchCV(estimator=xgb.XGBClassifier(objective = 'binary:
↳logistic',random_state=1,

↳use_label_encoder=False),

                                param_grid = param_grid,
                                scoring = 'accuracy',
                                verbose = 1,
                                n_jobs=-1,
                                cv = cv).fit(X,y)

print(optimal_params.best_params_,optimal_params.best_score_)

```

Fitting 5 folds for each of 729 candidates, totalling 3645 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed: 1.4s
[Parallel(n_jobs=-1)]: Done 176 tasks    | elapsed: 42.8s
[Parallel(n_jobs=-1)]: Done 426 tasks    | elapsed: 2.2min
[Parallel(n_jobs=-1)]: Done 776 tasks    | elapsed: 3.8min
[Parallel(n_jobs=-1)]: Done 1226 tasks   | elapsed: 6.2min
[Parallel(n_jobs=-1)]: Done 1776 tasks   | elapsed: 9.0min
[Parallel(n_jobs=-1)]: Done 2426 tasks   | elapsed: 12.5min
[Parallel(n_jobs=-1)]: Done 3176 tasks   | elapsed: 16.1min

{'gamma': 0.1, 'learning_rate': 0.2, 'max_depth': 6, 'n_estimators': 25,
'reg_lambda': 0, 'scale_pos_weight': 1.25} 0.9932027494193271

[Parallel(n_jobs=-1)]: Done 3645 out of 3645 | elapsed: 18.7min finished

```

```

[78]: m3 = xgb.XGBClassifier(objective = 'binary:logistic',random_state=1,gamma=0.
↳1,learning_rate = 0.2,max_depth=6,

                                n_estimators = 25,reg_lambda =
↳0,scale_pos_weight=1.25,use_label_encoder=False).fit(X,y)

```

0.10.4 Stacking Classifier

```

[79]: ensemble_model = StackingClassifier(estimators=[('dt',m1),('rf',m2),('xgb',m3)],

↳

↳final_estimator=LogisticRegression(random_state=1,max_iter=10000),n_jobs=-1,

                                cv =
↳StratifiedKfold(n_splits=5,shuffle=True,random_state=1))

ensemble_model.fit(X,y)

```

```

[79]: StackingClassifier(cv=StratifiedKfold(n_splits=5, random_state=1, shuffle=True),
                        estimators=[('dt',
                                    DecisionTreeClassifier(max_depth=6,
                                                            max_features=25,
                                                            max_leaf_nodes=15,
                                                            random_state=1)),

```

```

('rf',
 RandomForestClassifier(max_features=1,
                        n_estimators=500,
                        n_jobs=-1,
                        random_state=1)),
('xgb',
 XGBClassifier(base_score=0.5, booster='gbtree',
               callbacks=None,
               colsample_by...
               interaction_constraints='',
               learning_rate=0.2, max_bin=256,
               max_cat_to_onehot=4,
               max_delta_step=0, max_depth=6,
               max_leaves=0, min_child_weight=1,
               missing=nan,
               monotone_constraints='()',
               n_estimators=25, n_jobs=0,
               num_parallel_tree=1,
               predictor='auto', random_state=1,
               reg_alpha=0, reg_lambda=0, ...)],
 final_estimator=LogisticRegression(max_iter=10000,
                                     random_state=1),
 n_jobs=-1)

```

```

[80]: print(confusion_matrix_data(X,y,ensemble_model))
      print(confusion_matrix_data(Xtest,ytest,ensemble_model))

```

```

Accuracy = 99.33860003674444
Recall = 99.61538461538461
FPR = 0.914526908195568
FNR = 0.38461538461538464
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      2817.0      26.0
Actual 1       10.0     2590.0

```

```

Accuracy = 99.55240581872435
Recall = 99.84802431610942
FPR = 0.7326007326007326
FNR = 0.1519756838905775
Confusion matrix =
      Predicted 0 Predicted 1
Actual 0      1355.0      10.0
Actual 1       2.0     1314.0

```